

# 6.8300 Pset 4 Problem 2 writeup

Zhi Ren

March 22, 2025

## 1 Edited code snapshot

```
def camera_param_to_rays(c2w, intrinsics, H=128, W=128):
    """
    ---- Given the camera parameters, generate rays for each pixel.

    ---- Args:
    ----- c2w: [4, 4] camera-to-world transform matrix
    ----- intrinsics: [fx, fy, cx, cy] camera intrinsic parameters
    ----- H: Height of the image
    ----- W: Width of the image

    ---- Returns:
    ----- ray_origins: [H, W, 3] origin points for rays
    ----- ray_directions: [H, W, 3] direction vectors for rays
    ---- """
    # NOTE: This function should be the same as
    # in the volumetric rendering problem

    #####
    # YOUR CODE HERE #
    #####

    # Hint: Generate ray origins and directions
    # for each pixel in the image
    # 1. Create a meshgrid of pixel coordinates
    # 2. Convert pixel coordinates to
    # camera coordinates using intrinsics
    # 3. Transform camera coordinates
    # to world coordinates using c2w

    device = c2w.device
```

```

if not torch.is_tensor(intrinsics):
    intrinsics = torch.tensor(intrinsics,
                               device=device, dtype=torch.float32)
fx, fy, cx, cy = intrinsics

i = torch.arange(W, device=device,
dtype=torch.float32).view(1, W).expand(H, W)
j = torch.arange(H, device=device,
dtype=torch.float32).view(H, 1).expand(H, W)

dirs = torch.stack([(i - cx) / fx,
                    (j - cy) / fy,
                    torch.ones_like(i)], dim=-1)
# shape: [H, W, 3]

ray_directions = torch.einsum('hwc,dc->hwd',
dirs, c2w[:3, :3])

ray_directions = ray_directions
/ torch.norm(ray_directions, dim=-1, keepdim=True)

ray_origins = c2w[:3, 3].expand(H, W, 3)

return ray_origins, ray_directions

def sphere_tracing(
    ray_origins,
    ray_directions,
    model,
    t_near=0.0,
    t_far=3.0,
    max_iter=256,
    epsilon=1e-4,
):
    """
    ---- Perform sphere tracing to find the
    ---- intersection of rays with the implicit model.

    ---- Args:
    ---- ray_origins: [H, W, 3] origin points
    ---- for rays
    ---- ray_directions: [H, W, 3] direction
    ---- vectors for rays
    ---- model: Implicit model to compute the SDF
    ---- t_near: Near plane distance

```

```

----- t_far: Far plane distance
----- max_iter: Maximum number of iterations
----- for sphere tracing
----- epsilon: Distance threshold for stopping

---- Returns:
----- image: [H, W, 3] rendered image
----- """

device = ray_origins.device
H, W, _ = ray_origins.shape

image = torch.zeros(H, W, 3, device=device)

t = torch.full((H, W), t_near, device=device)

for i in range(max_iter):
    points = ray_origins + t.unsqueeze(-1)
    * ray_directions # [H, W, 3]
    points_flat = points.view(-1, 3)
    # Flatten to [N, 3] for the model

    # Evaluate SDF at these points.
    # model returns (sdf: [N,1], color: [N,3]),
    # but we only need sdf for marching.
    sdf, _ = model(points_flat)
    sdf = sdf.view(H, W) # Reshape to [H, W]

    active = (torch.abs(sdf) >= epsilon) & (t < t_far)

    if not active.any():
        break

    t[active] = t[active] + sdf[active]

final_points = ray_origins +
t.unsqueeze(-1) * ray_directions

hit_mask = (t < t_far) & (torch.abs(sdf) < epsilon)

if hit_mask.any():
    p_hit = final_points[hit_mask]
    # shape: [N_hit, 3]
    _, hit_colors = model(p_hit)
    # shape: [N_hit, 3]

```

```

        image[hit_mask] = hit_colors

    return image

def sample_points_on_rays(
    ray_origins, ray_directions,
    num_samples=64, t_near=0.0,
    t_far=3.0
):
    """
    ----Sample points along the rays.

    ----Args:
    -----ray_origins: [H, W, 3] ray origin points
    -----ray_directions: [H, W, 3] ray direction vectors
    -----num_samples: Number of sample points along each ray
    -----t_near: Near plane distance
    -----t_far: Far plane distance

    ----Returns:
    -----points: [H, W, num_samples, 3] sampled points in 3D space
    -----ts: [num_samples] distances along the rays
    -----"""
    # TODO: Implement this function
    # 1. Generate uniformly spaced samples along each ray
    # 2. Compute the 3D coordinates of each sample point
    # 3. Return an array of sample points
    # with shape [H, W, num_samples, 3]

    t_samples = torch.linspace(t_near, t_far,
                               num_samples, device=ray_origins.device)

    t_samples_expanded = t_samples.view(1, 1, num_samples, 1)

    # Compute the 3D coordinates of each sample point
    points = ray_origins.unsqueeze(-2) +
             ray_directions.unsqueeze(-2) * t_samples_expanded

    return points, t_samples

def volume_rendering(densities, colors, deltas):
    """
    ----Perform volume rendering to compute pixel
    ----colors from densities and colors.

```

```

---- Args:
----- densities: [H, W, num_samples, 1] - density
----- values-at-each-sample-point
----- colors: [H, W, num_samples, 3] - colors
----- at-each-sample-point
----- deltas: [num_samples] - intervals
----- between-adjacent-sample-points

---- Returns:
----- image: [H, W, 3] - rendered image
---- """

# TODO: Implement this function
# 1. Initialize accumulated color and transmittance
# 2. For each sample along the ray:
#     - Compute alpha from density and delta
#     - Update accumulated color and transmittance
# 3. Return the final rendered image
device = densities.device
H, W, num_samples, _ = densities.shape
sigma = densities[..., 0]

deltas = deltas.view(1, 1, num_samples)

alpha = 1 - torch.exp(-sigma * deltas) # [H, W, num_samples]

exp_term = torch.exp(-sigma * deltas) # [H, W, num_samples]

T_inclusive = torch.cumprod(exp_term, dim=2)
# [H, W, num_samples]

ones = torch.ones(H, W, 1, device=device)
T_exclusive = torch.cat([ones, T_inclusive[..., :-1]],
dim=2) # [H, W, num_samples]

weights = T_exclusive * alpha # [H, W, num_samples]

weights = weights.unsqueeze(-1)

image = torch.sum(weights * colors, dim=2) # [H, W, 3]

return image

```

## 2 Rendered images

Below are the rendered images from sphere tracing and volumetric rendering algorithms.




Figure 1: Rendered image from sphere tracing




Figure 2: Rendered image from volumetric rendering

## 3 Answer to questions

The reason why volumetric rendering is preferred to sphere tracing these days is mainly because volume-based methods are more robust. In volumetric rendering, the color is determined by summing over all the contributions along the marching ray, which is less prone to error if the geometry has very irregular shapes. In contrast, sphere tracing relies on finding the intersection point of the ray with the surface, which is more sensitive to the geometry of the scene.